

# Classes, debugging and optimisation

Paolo Joseph Baioni

March 26, 2025

## Exercise 1 - What you need to know

- ▶ The implemented matrix class is organized as **column-major**, i.e.  $A(i,j) = \text{data}[i + j * \text{rows()}]$ , conversion from 1d to 2d indexing is performed by the utility method `sub2ind`.
- ▶ Access to elements is implemented both in `const` and `non-const` versions, by overloading `operator()`.
- ▶ Data is private, *getter methods* expose what is needed to the user, both `const` and `non-const` versions are provided.
- ▶ Naive implementation of matrix-matrix multiplication is slow because it has low *data locality*, simply transposing the left matrix factor improves performance significantly<sup>1</sup>.
- ▶ The `#include <ctime>` header provides timing utilities, `tic()` and `toc(x)` macros start and stop the timer.

---

<sup>1</sup>See M. Kowarschik, C. Weiß. (2002). *Lecture Notes in Computer Science*. 213-232. DOI: 10.1007/3-540-36574-5\_10 for further details.

## Exercise 1.1

Starting from the provided implementation of the class for dense matrices (and column vectors represented as 1-column matrices) based on `std::vector`, implement the following methods:

- ▶ transpose:  $A = A^T$ .
- ▶ `operator*`: matrix-matrix and matrix-vector multiplication.

## Exercise 1.2

- ▶ Transpose the first factor in matrix multiplication before performing the product.
- ▶ Compare the execution speed with respect to the previous implementation.

## Exercise 1.2 - Details

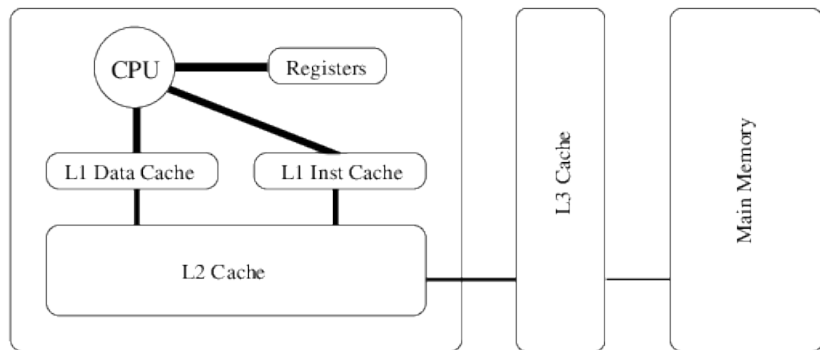


Figure: Typical memory layout of a computer.

## Exercise 1.2 - Details




Figure: Example with a row-major matrix.

## Exercise 1.2 - Details

---

1: double <i>sum</i> ;	1: double <i>sum</i> ;
2: double <i>a</i> [ <i>n</i> , <i>n</i> ];	2: double <i>a</i> [ <i>n</i> , <i>n</i> ];
3: <i>// Original loop nest:</i>	3: <i>// Interchanged loop nest:</i>
4: <b>for</b> <i>j</i> = 1 <b>to</b> <i>n</i> <b>do</b>	4: <b>for</b> <i>i</i> = 1 <b>to</b> <i>n</i> <b>do</b>
5: <b>for</b> <i>i</i> = 1 <b>to</b> <i>n</i> <b>do</b>	5: <b>for</b> <i>j</i> = 1 <b>to</b> <i>n</i> <b>do</b>
6: $sum += a[i, j];$	6: $sum += a[i, j];$
7: <b>end for</b>	7: <b>end for</b>
8: <b>end for</b>	8: <b>end for</b>

---

Figure: Example with a row-major matrix.

## Exercise 1.3

- ▶ Include the `Eigen/Dense` header.
- ▶ Use the `Eigen::Map` template class to wrap the matrix data and interpret it as `Eigen::MatrixXd`.
- ▶ Compare the execution speed with respect to the previous implementations.



# Bonus

- ▶ Link to Openblas
- ▶ Use the DGEMM function to perform the matrix-matrix product
- ▶ `matrix-0.4/matrix.{h,cpp}`

See [pacs-Labs/Labs/2019/03-cachealignment/matrix-0.4](https://pacs-labs.github.io/Labs/2019/03-cachealignment/matrix-0.4) for the solution

## Exercise 1.4

Going back to Ex. 1.2, perform the subsequent analysis:

- ▶ `coverage (lcov)`
- ▶ `memcheck (valgrind)`
- ▶ `profile (valgrind, kcachegrind)`

## Exercise 2

The program `integer-list` in the directory `02-bug` has:

- ▶ a compile error;
- ▶ a run-time error;
- ▶ a memory leak;
- ▶ bonus: a potential memory leak that is not captured by the `main`.

Find all the issues and fix them.

Get help by using `gdb` and `valgrind`.

The directory `02-bug-solution` contains the fixed code, please don't look at it before trying to solve the exercise by yourself!